

# Operating Systems

## CS220

Lecture 7

### Inter-Process Communication (IPC)

27th April 2021

By: Dr. Rana Asif Rehman

# What's in today's lecture

1. Process Concept
2. Process Manager Responsibilities
3. Process Scheduling
4. Operations on Processes
5. **Cooperating Processes**
6. **Inter-process Communication**

# Introduction

- A process has access to the memory which constitutes its own address space.
- When a child process is created, the only way to communicate between a parent and a child process is:
  - The parent receives the exit status of the child
- So far, we've discussed communication mechanisms only during process creation/termination
- Processes may need to communicate during their life time.


# Cooperating Processes

- **Independent** process cannot affect or be affected by the execution of another process.
- **Cooperating** process can affect or be affected by the execution of another process
- Advantages of process cooperation
  - Information sharing
  - Computation speed-up
  - Modularity
  - Convenience
- Dangers of process cooperation
  - Data corruption, deadlocks, increased complexity
  - Requires processes to synchronize their processing

# Inter-Process Communication (IPC)


- Mechanism for processes to communicate and to synchronize their actions.
- Two main types
  - Shared Memory
  - Message Passing

# Communications Models



(a)

**Message Passing**



(b)

**Shared Memory**

# Shared Memory (IPC)

- Uses producer-consumer paradigm
- Example of cooperating processes, *producer* process produces information that is consumed by a *consumer* process
- Implemented using buffers (bounded, non-bounded)
- How might you implement a bounded buffer scheme?

# Message Passing (IPC)

- **Message system** — processes communicate with each other **without** resorting to **shared variables**.
- IPC facility provides two operations:
  - **send**(message) — message size fixed or variable
  - **receive**(message)
- If P and Q wish to communicate, they need to:
  - establish a communication link between them
  - exchange messages via send/receive
- Implementation of communication link
  - physical (i.e., shared memory, hardware bus)
  - logical (direct/indirect, blocking/non-blocking, automatic/explicit buffering)



# Implementation Questions

- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?

# Message Passing Systems

- Exchange messages over a communication link
- Methods for implementing the communication link and primitives (send/receive):
  1. Direct or Indirect communications (Naming)
  2. Symmetric or Asymmetric communications (blocking versus non-blocking)
  3. Buffering

# Direct Communication

- Processes must **name each other** explicitly:
  - **send** (P, message) – send a message to process P
  - **receive**(Q, message) – receive a message from process Q
- Properties of communication link
  - Links are established automatically.
  - A link is associated with exactly one pair of communicating processes.
  - Between each pair there exists exactly one link.
  - Processes need to know each other's identity
  - The link may be unidirectional, but is usually bi-directional.

**Disadvantage:** a process must know the name or ID of the process(es) it wishes to communicate with

# Indirect Communication

- Messages are directed and received from mailboxes (also referred to as ports).
  - Each mailbox has a unique id.
  - Processes can communicate only if they share a mailbox.
- Properties of communication link
  - Link established only if processes have a shared mailbox
  - A link may be associated with many processes.
  - Each pair of processes may share several communication links.
  - Link may be unidirectional or bi-directional.

# Indirect Communication (Cont.)

- Operations provided by the OS
  - **create** a new **mailbox**
  - send and receive messages through mailbox
  - destroy a mailbox
- Primitives are defined as:
  - **send**(A, message) – send a message to mailbox A
  - **receive**(A, message) – receive a message from mailbox A

# Indirect Communication (Cont.)

- Mailbox sharing
  - P1, P2, and P3 share mailbox A.
  - P1, sends; P2 and P3 receive.
  - Who gets the message?
- Solutions
  - Allow a link to be associated with at most two processes.
  - Allow only one process at a time to execute a receive operation.
  - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

# Synchronization

- Message passing can be blocking or non-blocking
- Four types of synchronization
  - Blocking send:
    - sender blocked until message received by mailbox or process
  - Nonblocking send:
    - sender resumes operation immediately after sending
  - Blocking receive:
    - receiver blocks until a message is available
  - Nonblocking receive:
    - receiver returns immediately with either a valid or null message.

# Buffering

- Messages exchanged by processes reside in temporary queue
- Three ways to implement queues

1. Zero capacity

No messages may be queued within the link, requires sender to block until receiver retrieves message.

2. Bounded capacity

Link has finite number of message buffers. If no buffers are available then sender must block until one is freed up.

3. Unbounded capacity

Link has unlimited buffer space, consequently send never needs to block.



# Communication in Client/Server Systems

- Sockets
- Remote procedure calls
- Pipes
- What types of Inter-Process Communication might each of these use?

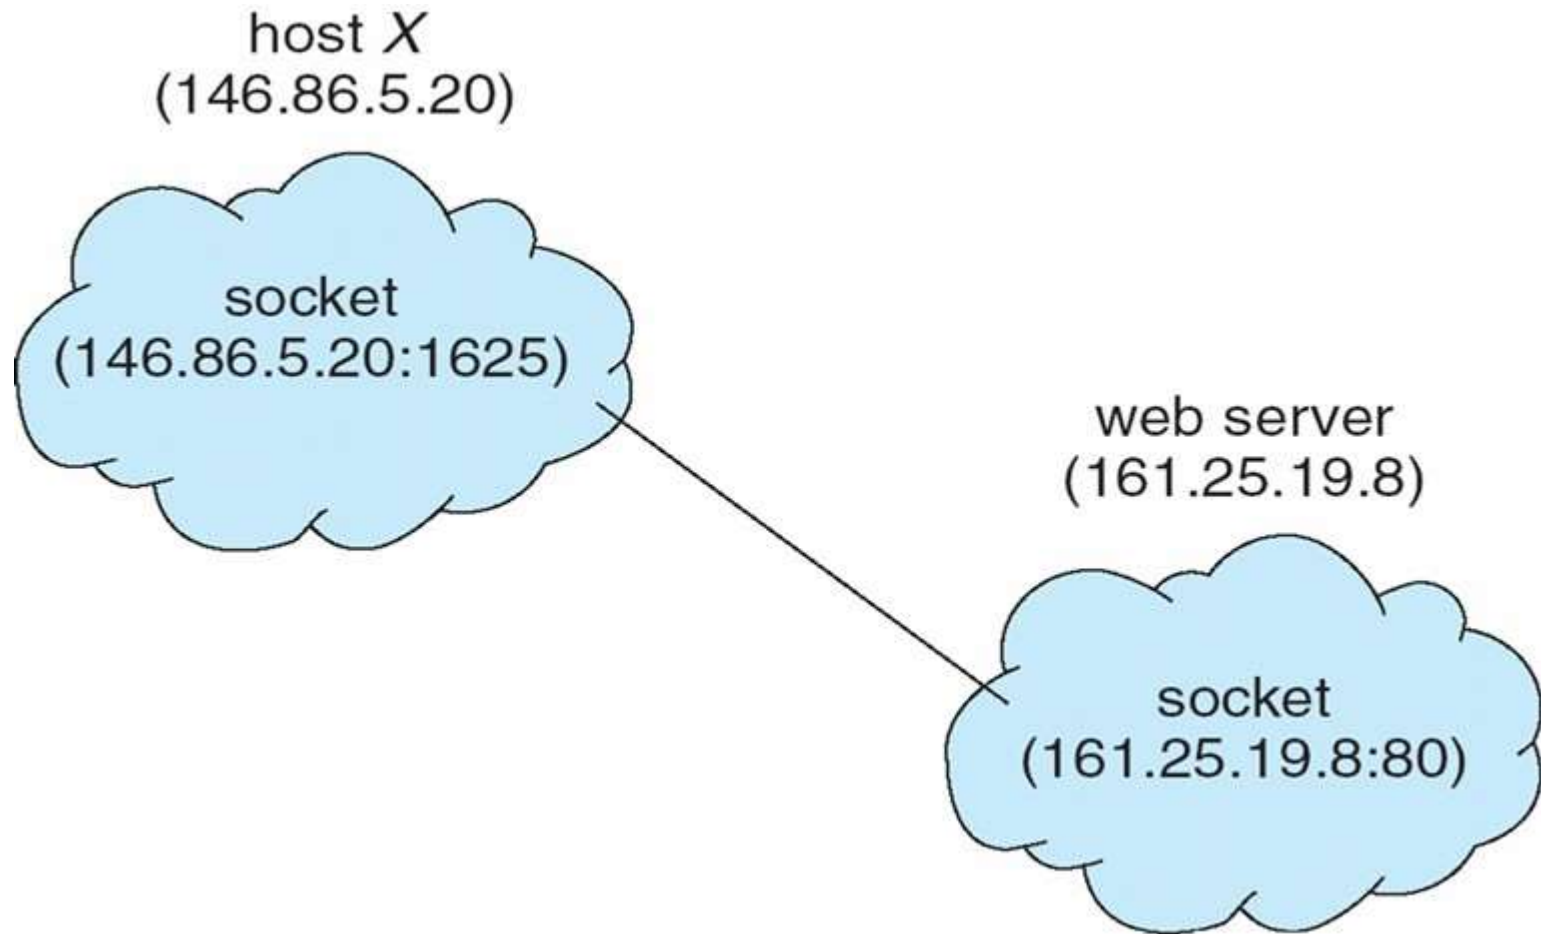
# 1. Sockets

- Defined as a point of communication
- A pair of process communicating over a network employ a pair of sockets
- Socket = IP Address : Port Number
- All ports below 1024 are considered well known
  - FTP: 21
  - Telnet: 23
  - Web server: 80

# Sockets

- Client initiates a requests for connection, it is assigned a port by host computer
- Connection on one host must be unique
- Sockets can be
  - Connection-oriented (TCP) socket
  - Connectionless (UDP) socket
- Considered Low level mode of communication
  - Shares unstructured stream of bytes

# Sockets



## 2. Remote Procedure Calls (RPC)

- Remote procedure call (RPC) abstracts procedure calls between processes on networked systems.
- **Stubs** – client-side proxy for the actual procedure on the server.
- The client-side stub locates the server and marshals the parameters.
- The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server.

# RPC: An Example

**Client Program:**

```
sum = server->Add(3,4);
```

**Client Stub:**

```
Int Add(int x, int y) {  
    Alloc message buffer;  
    Mark as "Add" call;  
    Store x, y into buffer;  
    Send message;  
}
```

**RPC Runtime:**

```
Send message to server;
```

**Server Program:**

```
int Add(int x, int, y) {}
```

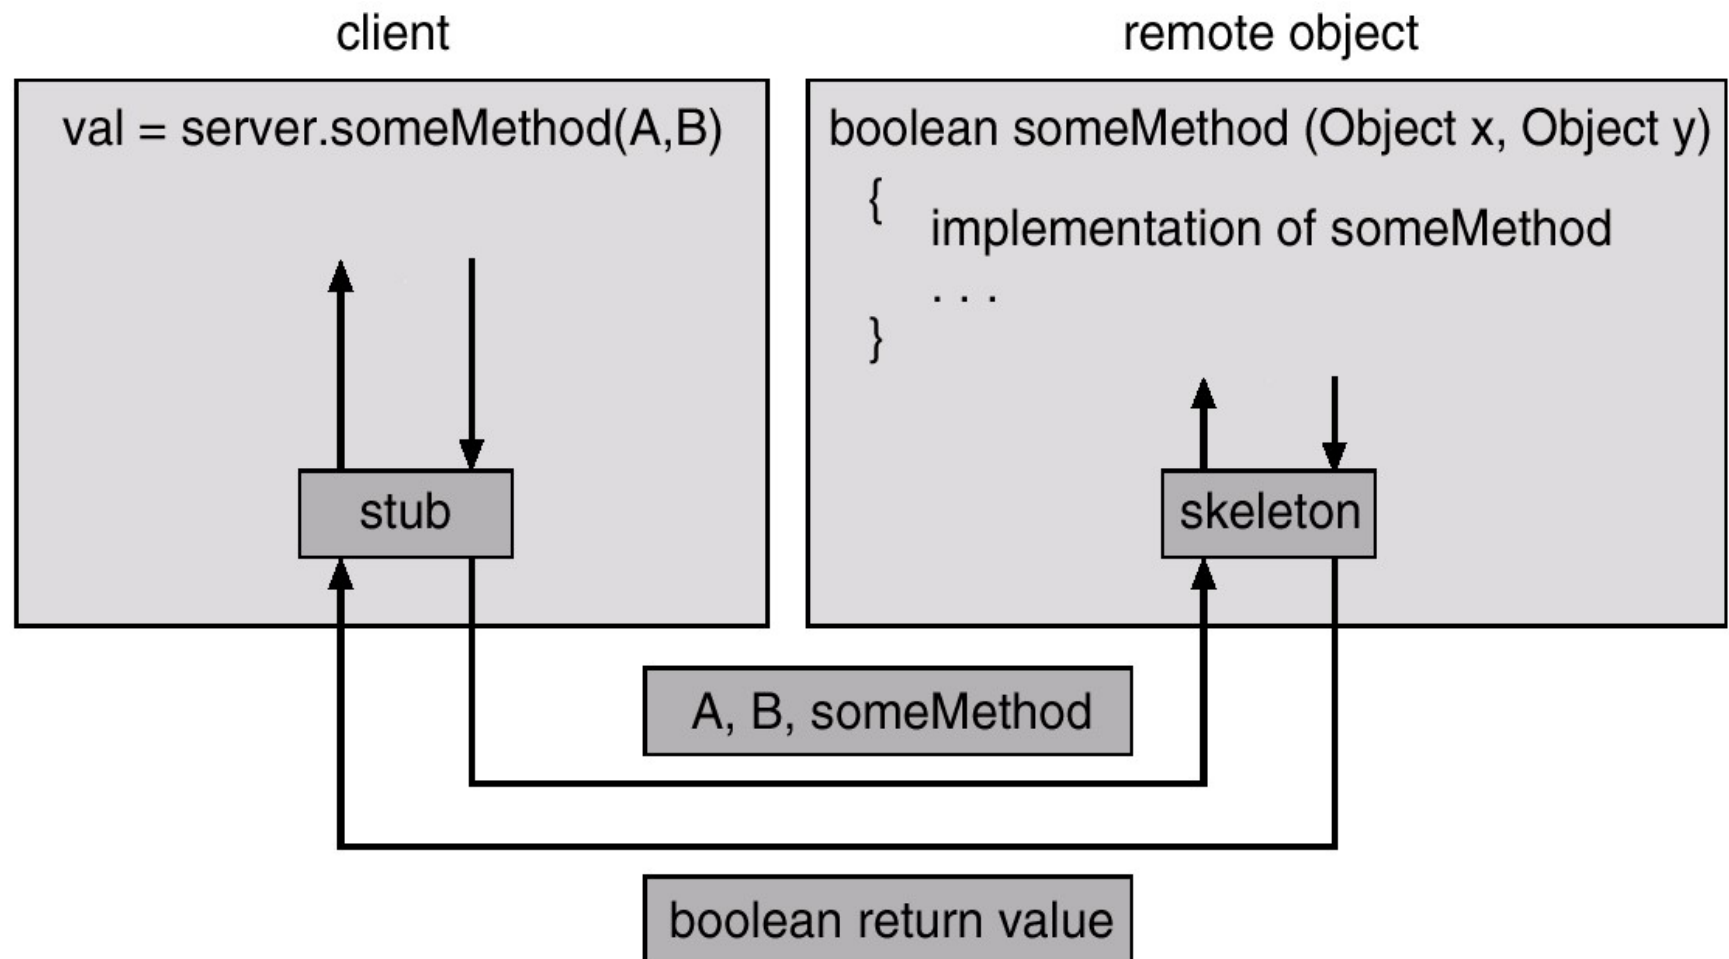
**Server Stub:**

```
Add_Stub(Message) {  
    Remove x, y from buffer  
    r = Add(x, y);  
}
```

**RPC Runtime:**

```
Receive message;  
Dispatch, call Add_Stub;
```

# Marshalling Parameters

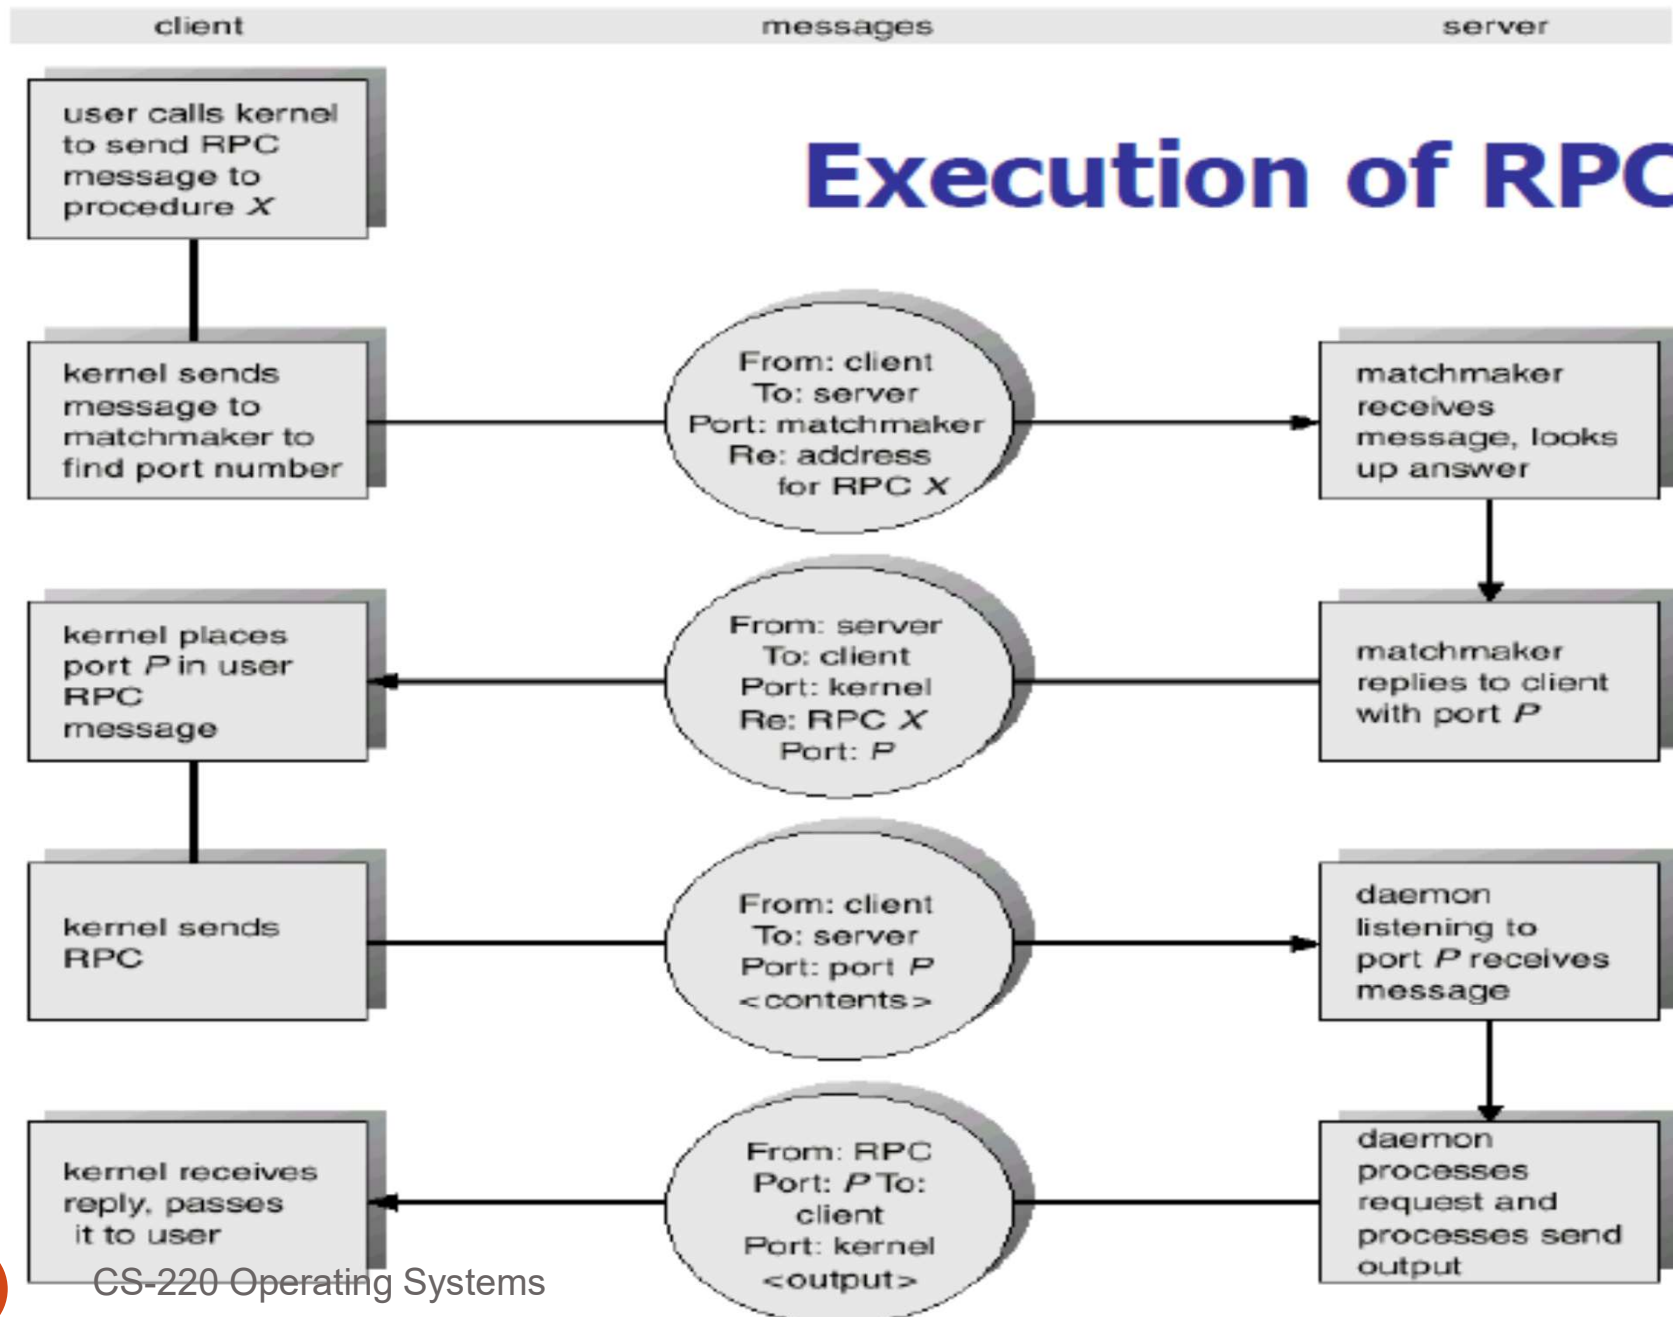


# Remote Procedure Calls

- Problem exists in data representation
  - Big endian
  - Little endian
- Many RPC systems define a machine independent representation
  - External Data Representation (XDR)
- Semantics of a call
  - Messages are acted on *exactly once* rather than *at-most once*
- Port binding of RPC server
  - Fixed at compile time
  - Done dynamically by **rendezvous** mechanism (*match maker*)




# Execution of RPC



# 3. Pipes

- IPC mechanism in early UNIX systems
- Four issues must be considered
  - Unidirectional or bi-directional communication
  - If bidirectional then is it half duplex or full duplex
  - Must a relationship (such as *parent-child*) exist between the communicating processes?
  - Can pipe communicate over a network?

# Pipes: Shared info in kernel's memory



# Pipe Creation (Unix)

```
#include <unistd.h>
```

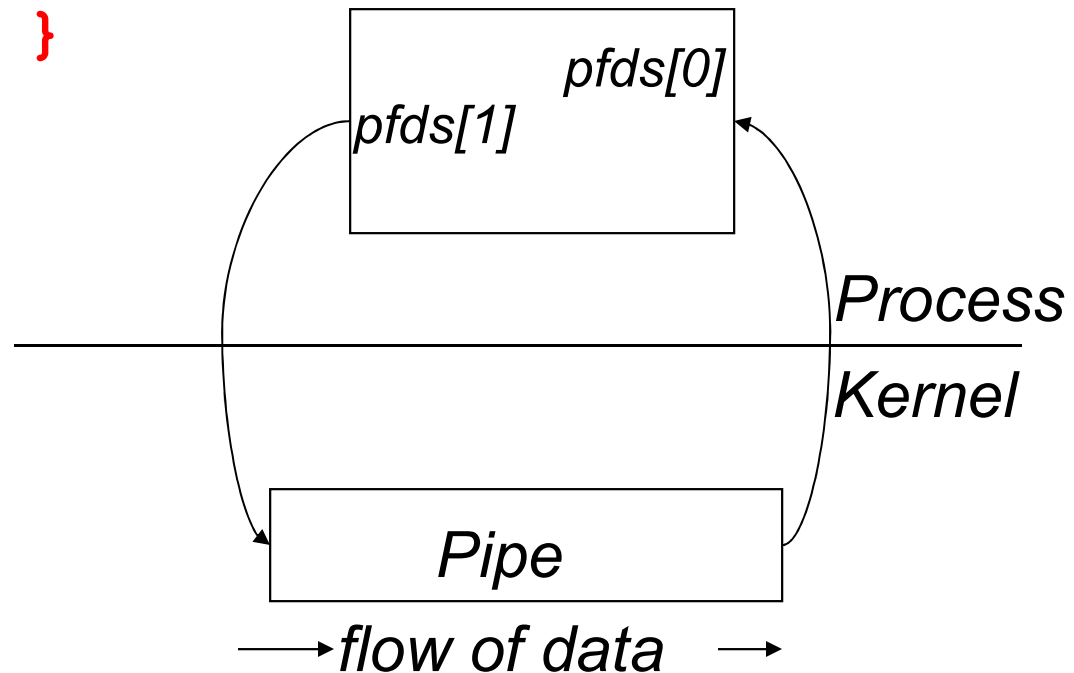
```
int pipe(int filedes[2]);
```

- Creates a pair of file descriptors pointing to a pipe inode
- Places them in the array pointed to by filedes
  - `filedes[0]` is for reading
  - `filedes[1]` is for writing.
- On success, zero is returned.
- On error, -1 is returned

# Pipe Creation

```
int main()
{  int pfd[2];

  if (pipe(pfd) == -1)
    {  perror("pipe");
        exit(1);  }
}
```



# Reading/Writing from/to a Pipe

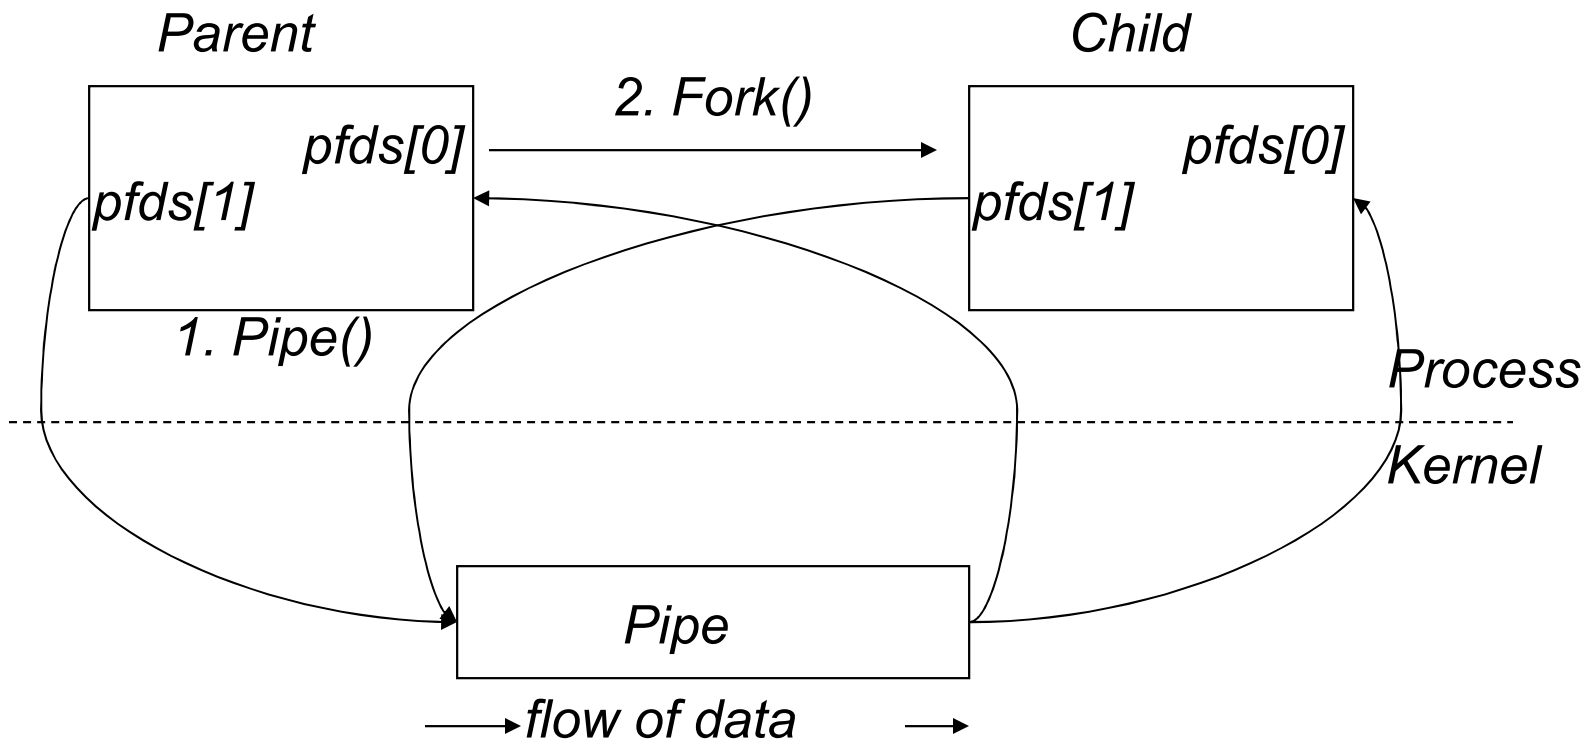
- `int read(int filedescriptor, char *buffer, int bytetoread);`
- `int write(int filedescriptor, char *buffer, int bytetowrite);`

# Example

```
int main()
{ int pfds[2];
  char buf[30];
  if (pipe(pfds) == -1) {
      perror("pipe");
      exit(1); }
  printf("writing to file descriptor #%d\n", pfds[1]);
  write(1, "test", 5);?

  printf("reading from file descriptor
  #0\n");
  read(0, buf, 5);?
  printf("read %s\n", buf);
}
```

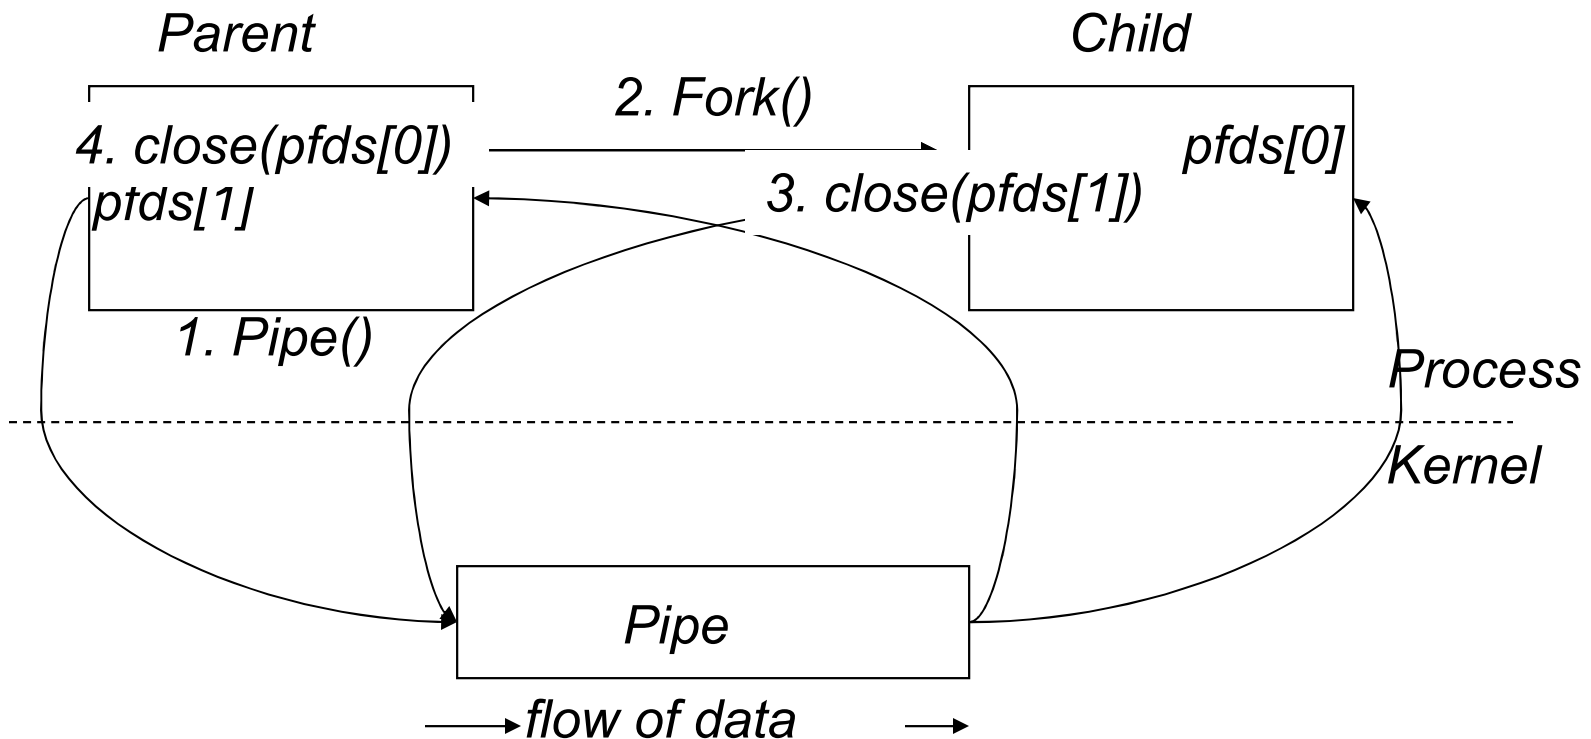
# A Channel between two processes





# A Channel between two processes

- To allow one way communication each process should close one end of the pipe.



# An Example of pipes with fork

```
int main() {  
    int pfds[2];  
    char buf[30];  
    pipe(pfds);.....1  
    if (fork()==0) .....2  
    { close(pfds[0]);.....3  
      printf(" CHILD: writing to the pipe\n");  
      write(pfds[1], "test", 5);  
      printf(" CHILD: exiting\n");  
      exit(0);  
    }  
    else { close(pfds[1]);.....4  
          printf("PARENT: reading from pipe\n");  
          read(pfds[0], buf, 5);  
          printf( "PARENT: read \"%s\"\n", buf);  
          wait(NULL);  
        }  
}
```


# Pipes

- Two types of pipes
  - Ordinary pipes
  - Named pipes

# Ordinary pipes

- **Ordinary pipes** allow communication in standard producer-consumer form
- **Producer** writes to one end of the pipe (the *write-end*)
- **Consumer** reads from other end of the pipe (the *read-end* of the pipe)
- Ordinary pipes are therefore unidirectional
- Required parent-child relationship between communicating processes
- In windows known as *Anonymous pipes*

# Ordinary pipe



# Named pipes

- Named pipes are more powerful than ordinary pipes
- Communication is bidirectional
- No parent-child relationship is necessary between the communicating processes
- Several processes can use the named pipe for communication
- Provided on both UNIX and Windows
- Called FIFO in UNIX

# References

- Operating System Concepts (Silberschatz, 9<sup>th</sup> edition)  
Chapter 3