

Operating Systems

CS220

Lecture 8

Threads

3rd May 2021

By: Dr. Rana Asif Rehman

What's in today's lecture

- Thread Concept
- Multithreading Models
- User & Kernel Threads
- Pthreads
- Threads in Solaris, Linux, Windows

Motivation

- Most modern applications are multithreaded
- Threads run within application
- Multiple tasks with the application can be implemented by separate threads
 - Update display
 - Fetch data
 - Spell checking
 - Answer a network request
- Process creation is heavy-weight while thread creation is light-weight
- Can simplify code, increase efficiency
- Kernels are generally multithreaded

Introduction

- Each process has
 1. Own Address Space
 2. Single thread of control
- A process model has two concepts:
 1. Resource grouping
 2. Execution
- Sometimes it is useful to separate them

Unit of Resource Ownership

- A process has an
 - Address space
 - Open files
 - Child processes
 - Accounting information
 - Signal handlers
- If these are put together in a form of a process, can be managed more easily

Unit of Dispatching

- Path of execution
 - Program counter: which instruction is running
 - Registers:
 - holds current working variables
 - Stack:
 - Contains the execution history, with one entry for each procedure called but not yet returned
 - State
- Processes are used to group resources together
- Threads are the entities scheduled for execution on the CPU
- Threads are also called *lightweight* process (LWP)

Its better to distinguish between the two concepts

Heavy weight process

Address
space/Global
Variables
Open files
Child processes
Accounting info
Signal handlers
Program counter
Registers
Stack
State

In case of multiple
through per process

Unit of Resource

Split

Unit of Dispatcher

Address space/Global
Variables
Open files
Child processes
Accounting Info
Signal Handlers

Program Counter
Reg
Stac
Stat

Program Counter
Registers
S
S

Program Counter
Registers
Stack
State

Share

Light weight processes

Process Vs. Thread


- Process

- All resources allocated: IPC channels, files etc..
- A virtual address space that holds the process image
- Protected access to processors, other processes, I/O resources, and files


- Threads

- a dispatchable unit of work
- an execution state: running, ready, etc..
- saved thread context (when not running)
- an execution context: PC, SP, other registers
- a per-thread stack
- Access to the memory and resources of the process it belongs to
 - all threads of the same process share this

Process Vs. Threads




(a) Three threads, each running in a separate address space



(b) Three threads, sharing the same address space

The Thread Model




Each thread has its own stack

Multicore Programming


- **Multicore** or **multiprocessor** systems putting pressure on programmers, challenges include:
 - Dividing activities
 - Balance
 - Data splitting
 - Data dependency
 - Testing and debugging
- **Parallelism** implies a system can perform more than one task simultaneously
- **Concurrency** supports more than one task making progress
 - Single processor / core, scheduler providing concurrency

Concurrency vs. Parallelism

- Concurrent execution on single-core system:



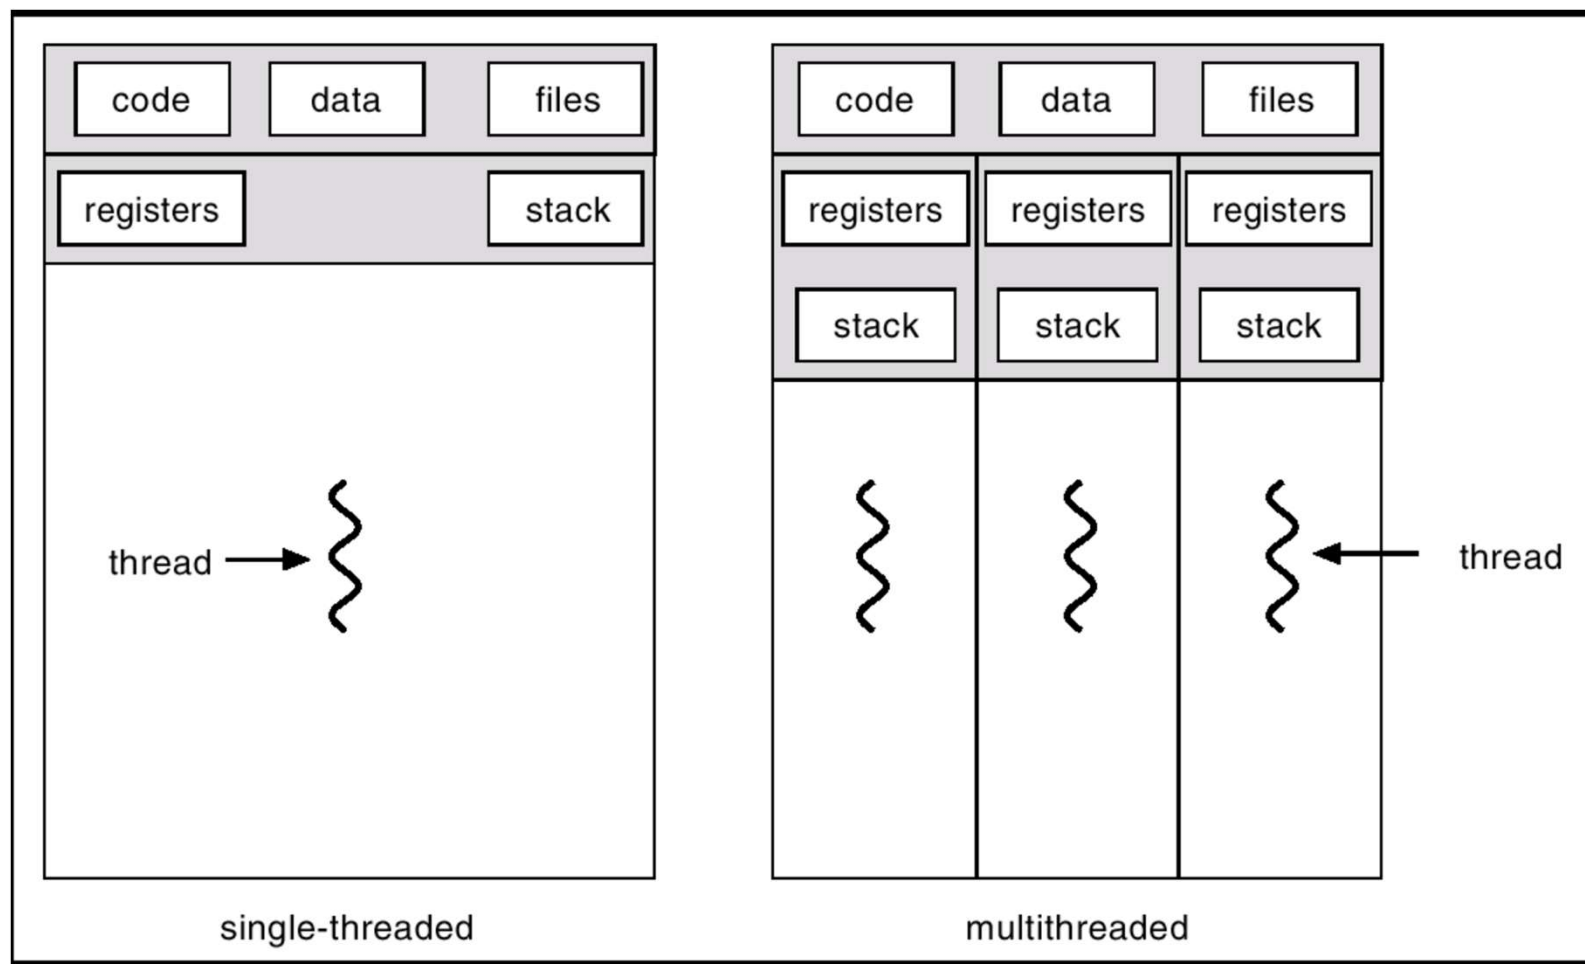
- Parallelism on a multi-core system:



Multicore Programming (Cont.)

- Types of parallelism
 - **Data parallelism** – distributes subsets of the same data across multiple cores, same operation on each
 - **Task parallelism** – distributing threads across cores, each thread performing unique operation
- As # of threads grows, so does architectural support for threading
 - CPUs have cores as well as *hardware threads*
 - Consider Oracle SPARC T4 with 8 cores, and 8 hardware threads per core

Single and Multithreaded Processes






Figure 4.1 Threads and Processes [ANDE97]

Threads

- Allow multiple execution paths in the same process environment
(Within an address space, we can have more units of execution: **threads**)
- All the threads of a process share the same address space and the same resources
- But have own set of Program counter, Stack etc
- The first thread starts execution with
 - **`int main(int argc, char *argv[])`**
- The threads appear to the Scheduling part of an OS just like any other process

Threads

- **Advantages**

- operations on threads (creation, termination, scheduling, etc..) are **cheaper** than the corresponding operations on processes
- inter-thread communication is supported through **shared memory** without kernel intervention
- Responsiveness, Resource Sharing, Utilization of MP Architectures


- **Disadvantages**

- easy to introduce race conditions
- synchronization is necessary

Thread Usage


- Less time to create a new thread than a process
 - the newly created thread uses the current process address space
 - no resources attached to them
- Less time to terminate a thread than a process.
- Less time to switch between two threads within the same process, because the newly created thread uses the current process address space.
- Less communication overheads
 - threads share everything: address space, in particular. So, data produced by one thread is immediately available to all the other threads
- Performance gain
 - Substantial Computing and Substantial Input/Output
- Useful on systems with multiple processors

1. Single threaded database server



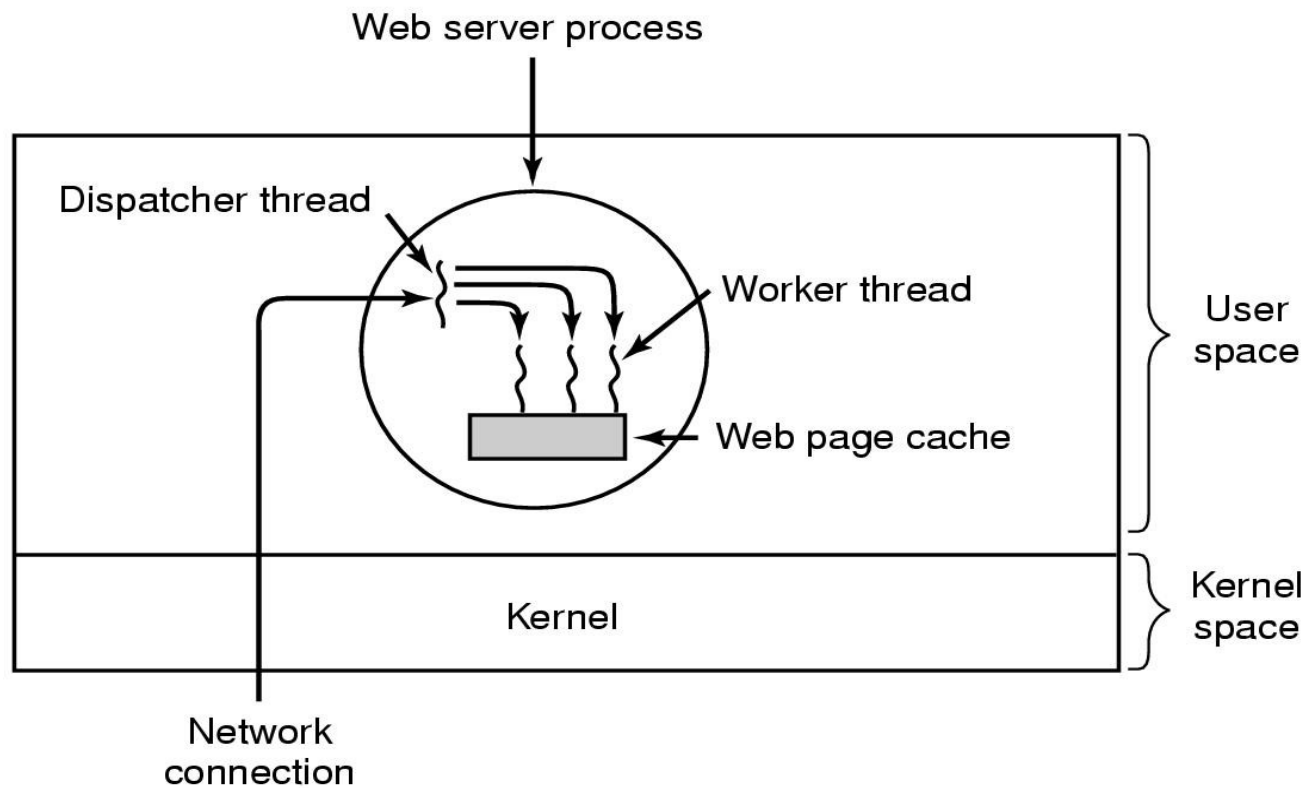
- Handles multiple clients
 - Either handle the requests sequentially
 - Or multiplex explicitly by creating multiple processes
- Problems
 - Unfair for quick requests, occurring behind lengthy request
 - Complex and error prone
 - Heavy IPC required

1. Multithreaded database server



- Assign a separate thread to each request
- As fair as in the multiplexed approach.
- The code is as simple as in the sequential approach, since the address space is shared - all variables are available
- Some synchronization of access to the database is required, this is not terribly complicated.

e.g. A multithreaded web server




2. Background Processing


- Consider writing a GUI-based application that uses:
 - Mouse
 - Keyboard input
 - Handles various clock-based events
- In a single threaded application, if the application is busy with one activity, it cannot respond (quickly enough) to other events, such as mouse or keyboard input.
- Handling such concurrency is difficult and complex
- But simple in a multithreaded process



e.g. A word processor with 3 threads



3. Parallel Algorithms e.g. Merge Sort



- Sort some data items on a shared-memory parallel computer.
- Our task is merely to implement a multithreaded sorting algorithm.
 - Divide the data into four pieces
 - Have each processor sort a different piece.
 - Two processors each merge two pieces
 - One processor merges the final two combined pieces.

Threads

- User Threads
- Kernel Threads

User Threads

- All thread management is done by the application. A user thread maintains all its state in user space.
- Thread switching does not require kernel mode privileges (no mode switch)
- Scheduling is application specific
- **The kernel is not aware of the existence of user threads**
- **Examples**
 - POSIX *Pthreads*
 - Mach *C-threads*
 - Solaris *threads*

User thread (cont..)

- Threads library contains code for:
 - creating and destroying threads
 - passing messages and data between threads
 - scheduling thread execution
 - saving and restoring thread contexts

Kernel Threads

- Kernel threads are supported directly by the operating system
- The kernel performs creation, scheduling, and management
- Kernel threads are generally slower to create and manage than user threads
- However since the kernel manages the threads, if a thread performs a system call, the kernel can schedule another thread
- OS Kernel maintains context information for the process and the threads (LWP)
- Scheduling is done on a thread basis
- Examples
 - Windows NT/XP/Vista/7, Solaris, Tru64 UNIX, Linux

Thread Implementation

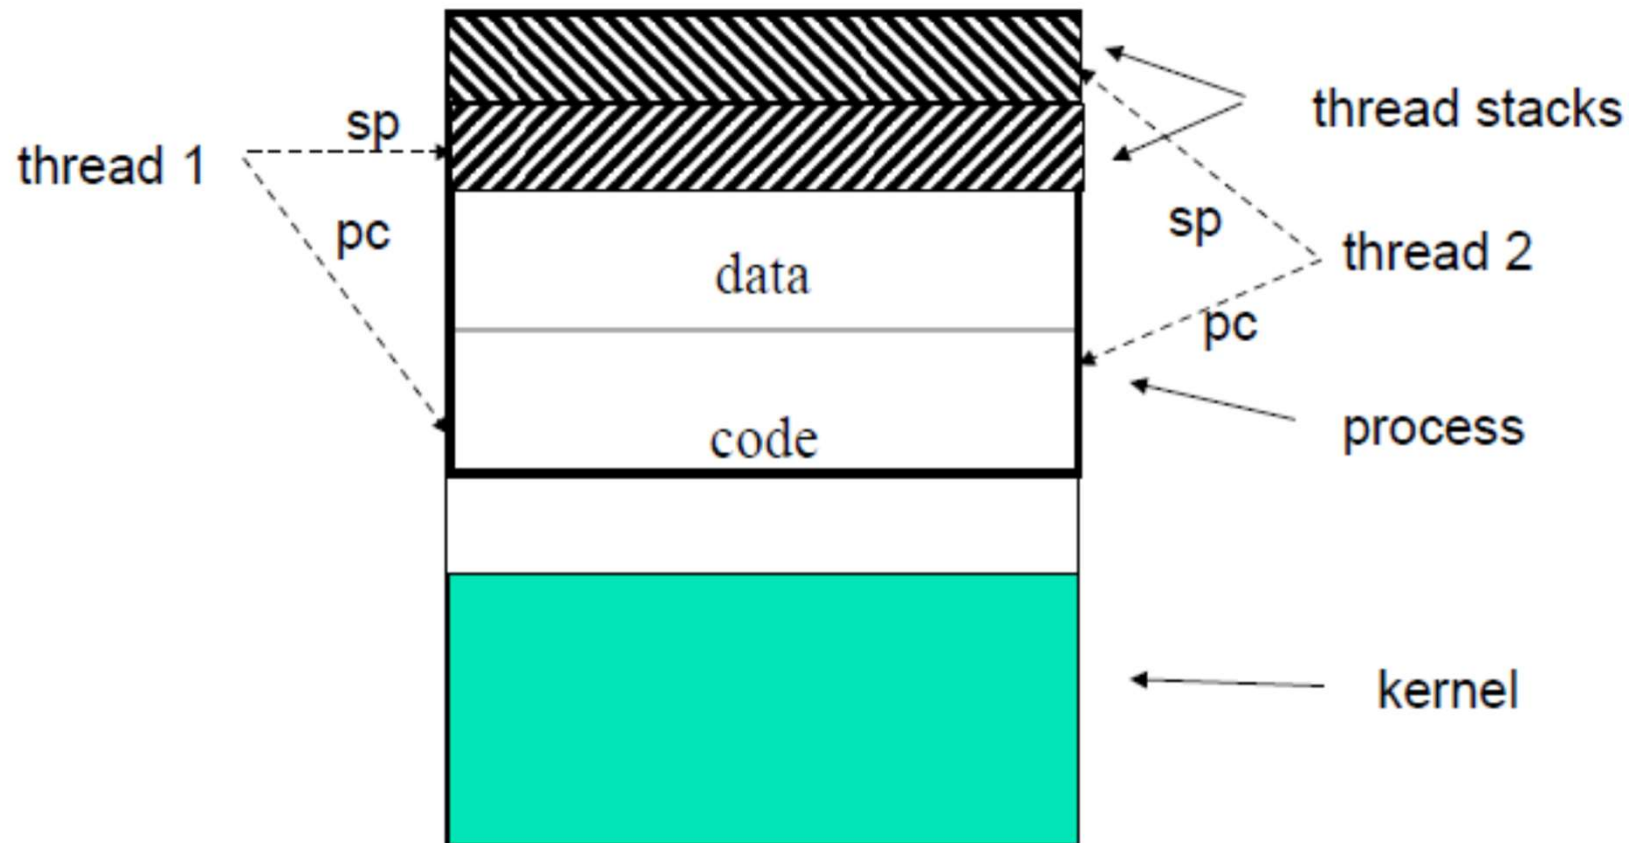
- User-level threads

- Implemented as a thread library, which contains the code for thread creation, termination, scheduling, and switching
- Kernel sees one process and it is unaware of its thread activity
- can be preemptive or not (co-routines)

- Kernel-level threads

- Thread management done by the kernel

User-Level Thread Implementation



User-Level vs. Kernel-Level Threads

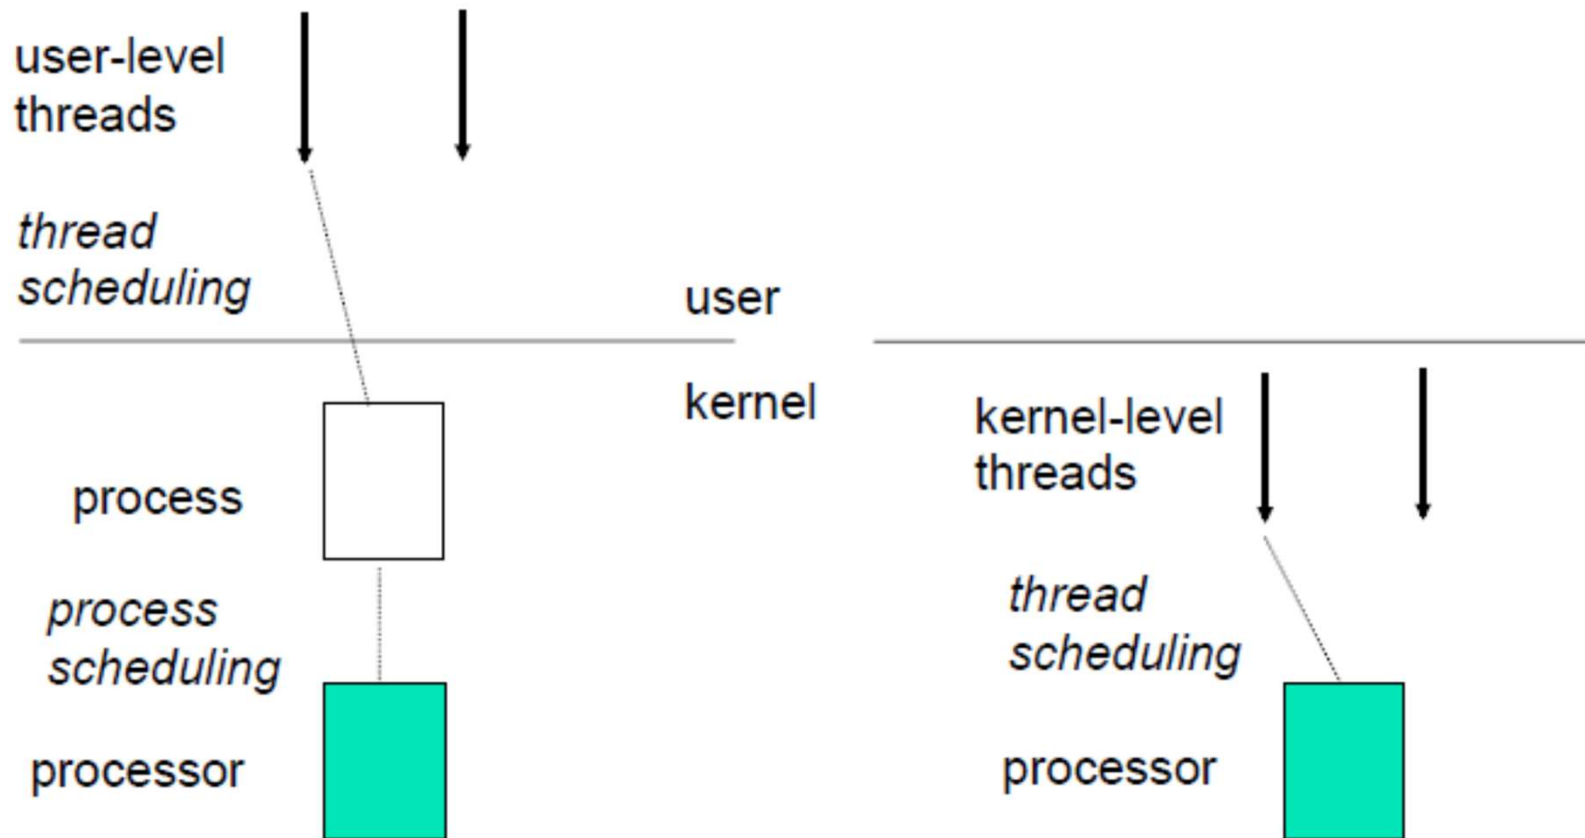
- Advantages of the user-level threads

- Performance: low-cost thread operations (do not require crossing protection domains), fast context switching
- Flexibility: scheduling can be application specific
- Portability: user-level thread library easy to port

- Disadvantages of the user-level threads

- If a user-level thread is blocked in the kernel, the entire process (all threads of that process) are blocked
- cannot take advantage of multiprocessing (the kernel assigns one process to only one processor)

User-Level vs. Kernel-Level Threads




Multithreading Models

- Many-to-One
- One-to-One
- Many-to-Many


Many-to-One Model

- Many user-level threads mapped to single kernel thread
- Used on systems that do not support multiple kernel threads
- E.g., Solaris green threads; GNU Portable Threads




One-to-One Model

- Each user-level thread maps to kernel thread.
- Examples
 - Windows NT/XP/Vista/7
 - Linux, Solaris 9



Many-to-Many Model

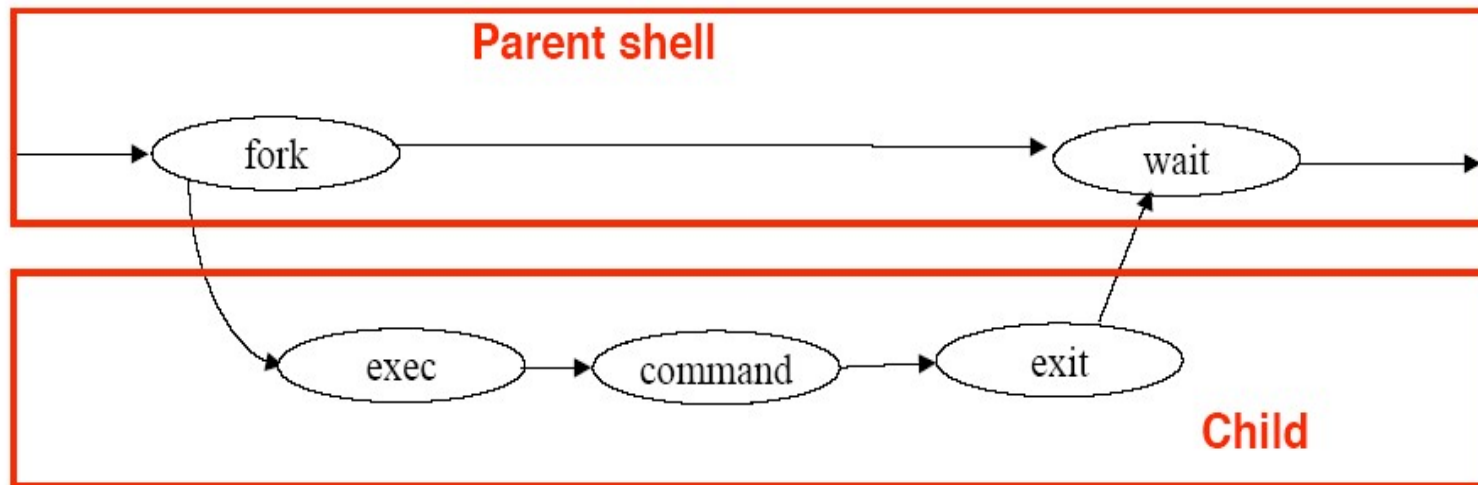
- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Solaris 2
- HP-UX, Tru64 UNIX



Multithreaded Programming Issues – fork and exec system calls

- When a thread (associated with process A) calls fork, two things can happen:
 - The new process duplicates all threads associated with process A
 - The new process will be single-threaded
- Some Unix operating systems support these two versions of fork
- Typically, the exec system call is used after a fork system call
- The exec system call:
 - Loads a binary file into memory
 - Destroys the memory image containing the exec system call
 - Starts its execution

How the unix shell runs commands



- when you type a command, the shell forks a clone of itself
- the child process makes an exec call, which causes it to stop executing the shell and start executing your command
- the parent process, still running the shell, waits for the child to terminate

Multithreaded Programming Issues – Cancellation

- Thread cancellation is the task of terminating a thread before it is completed
 - For example: assume that multiple threads are searching a database. As soon as one thread returns the search result, we can terminate the remaining threads
- A thread to be cancelled is referred to as a target thread
- Thread cancellation can happen in two ways:
 - **Asynchronous cancellation:** One thread immediately terminates the target thread
 - **Deferred cancellation:** the target thread can periodically checks if it should terminate. This allows the cancellation to happen in an orderly manner

Multithreaded Programming Issues – Cancellation

- Thread cancellation is not as easy as it appears
 - What about resources allocated to a thread
 - A thread might be cancelled while in the middle of updating a shared variable
 - This becomes especially troublesome with asynchronous cancellation
 - An OS usually reclaim all system resources from a cancelled thread. But often does not reclaim all resources. Why?
- Deferred cancellation happens when a thread can be safely cancelled. This is referred to as **cancellation points**
- The Pthreads API provides **cancellation points**

Multithreaded Programming Issues – Signal Handling

- A signal is used to notify a process that a particular event has occurred
 - Examples of signals are illegal memory access, division by zero, etc
- A signal might occur synchronously and asynchronously
- All signals follow the same pattern
 - A signal is generated by the occurrence of a particular event
 - A generated signal is delivered to a process
 - Once delivered, the signal must be dealt with
- **Synchronous signals** are generated by events internal to a running process
- **Asynchronous signals** are generated by events external to a running process

Multithreaded Programming Issues – Signal Handling

- **Synchronous signals:**
 - Generated by events internal to the running process
 - Examples include illegal memory access, division by zero, etc
 - Synchronous signals are delivered to the same process that performed the operation causing the signal
- **Asynchronous signals:**
 - Generated by events external to the running process
 - Examples include terminating a process by specific keystrokes (e.g., control c)
 - Asynchronous signals are more complicated

Multithreaded Programming Issues

– Signal Handling

- Every signal, whether synchronous or asynchronous, is handled in two ways
 - A default signal handler
 - A user-defined signal handler
- By default, every signal has a **default signal handler** that is run by the kernel
- This default signal handler can be overwritten by the user-defined signal handler
- **Single-threaded programs:** straightforward, signals are always delivered to the process
- **Multithreaded programs:** more complicated

Multithreaded Programming Issues – Signal Handling

- When a signal is delivered to a multithreaded program, the following can happen:
 - Deliver the signal to the thread to which the signal applies
 - Deliver the signal to every thread in the process
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process
- **Examples:**
 - A terminating signal should be sent to all thread in the process
 - Solaris 2 implements the fourth option (i.e., creates a special thread within each process solely for signal handling)

Multithreaded Programming Issues

- Thread Pools

- Creating threads can be time consuming
- Too many threads can bog down the system
- Thread pools help with this problem
- Threads are pre-allocated
- The number of threads available at a given time is fixed
- Some systems may adjust the thread pool size depending on usage

Multithreaded Programming Issues

- Thread Specific Data/ Thread-local Storage (TLS)
 - Threads belonging to the same process share the process data. This provided the benefit of multithreaded programming
 - However, in some instances, each thread might need its own specific data.
 - For example, a transaction processing multithreaded application might service each transaction in a separate thread
 - Most thread libraries such as Win32 and Pthreads provides support for thread specific data

Thread Implementation (POSIX)

- The most important of these APIs, in the Unix world, is the one developed by the group known as POSIX.
- POSIX is a standard API supported
- Portable across most UNIX platforms.
- PTHREAD library contains implementation of POSIX standard
- To link this library to your program use *-lpthread*
 - **gcc MyThreads.c -o MyThreadExecutable -lpthread**

POSIX Thread (Pthread)

- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization.
- API *specifies* behavior of the thread library, *implementation* is up to developer of the library.
- Common in UNIX operating systems.

Pthread API

- **thread creation and termination**

- *pthread_create(&tid, NULL, start_fn, arg);*
- *pthread_exit(status);*

- **thread join**

- *pthread_join(tid, &status);*

- **mutual exclusion**

- *pthread_mutex_lock(&lock);*
- *pthread_mutex_unlock(&lock);*

- **condition variable**

- *pthread_cond_wait(&c, &lock);*
- *pthread_cond_signal(&c);*

Condition Variables (Example)

- thread 1

```
pthread_mutex_lock(&lock);  
while (!my-condition)  
    pthread_cond_wait(&c, &lock);  
do_critical_section();  
pthread_mutex_unlock(&lock);
```

- thread 2


```
pthread_mutex_lock(&lock);  
my-condition = true;  
pthread_mutex_unlock(&lock);  
pthread_cond_signal(&c);
```

```


#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5
void * PrintHello(void *threadid)
{
    printf("\n%d: Hello World!\n", threadid);
    pthread_exit(NULL);
}
int main()
{
    pthread_t threads[NUM_THREADS]; /* the thread identifier */
    int rc, t;
    for (t = 0; t < NUM_THREADS; t++)
    {
        printf("Creating thread %d\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *)&t);
        if (rc)
        {
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }
    pthread_exit(NULL);
}

```

Solaris 2 Threads



Solaris Process



Linux Threads

- Linux refers to them as tasks rather than threads
- Thread creation is done through clone() system call
- Clone() allows a child task to share the address space of the parent task (process)

Windows NT Threads

- Implements the one-to-one mapping.
- Each thread contains
 - a thread ID
 - register set
 - separate user and kernel stacks
 - private data storage area

References

- Chapter 2, Modern Operating Systems, Tanenbaum
- Chapter 4, Silberschartz Operating System Concept 9th Edition.
- <http://www.thegeekstuff.com/2012/04/create-threads-in-linux/>
- <http://www.yolinux.com/TUTORIALS/LinuxTutorialPosixThreads.html>
- http://pubs.opengroup.org/onlinepubs/7908799/xsh/pthread_create.html