

Operating Systems

CS220

Lecture 13

Process Synchronization-II

21th June 2021

By: **Dr. Rana Asif Rehman**

Process Synchronization

Objectives

- The Critical-Section Problem
- Synchronization Hardware
- Semaphores
- Classical Problems of Synchronization

Synchronization Hardware

- Many systems provide hardware support for critical section code
- Uniprocessors – could disable interrupts
 - Currently running code would execute without preemption
 - Generally too inefficient on multiprocessor systems
 - Operating systems using this not broadly scalable
- Modern machines provide special atomic hardware instructions
 - Atomic = non-interruptable
 - Either test memory word and set value
 - Or swap contents of two memory words

Solution using TestAndSet

- Shared boolean variable **lock**., initialized to **false**

```
while (true) {  
    while ( TestAndSet (&lock ))  
        ; // do nothing  
        // critical section  
    lock = FALSE;  
    // remainder section  
}
```

```
int  
TestAndSet (boolean  
*lockValue)  
{  
    boolean rv;  
    rv = *lockValue;  
    *lockValue = true;  
    return rv;  
}
```

Solution using Swap

- Shared Boolean variable **lock** initialized to **FALSE**; Each process has a local Boolean variable **key**.
- Solution:

```
while (true) {  
    key = TRUE;  
    while ( key == TRUE)  
        Swap (&lock, &key );  
    // critical section  
    lock = FALSE;  
    // remainder section  
}
```

Definition:

```
void Swap (boolean *a, boolean *b)  
{  
    boolean temp = *a;  
    *a = *b;  
    *b = temp;  
}
```

Test-and-Set Instruction

- Mutual exclusion is assured: if P_i enters CS, the other processes are *busy waiting*
- Satisfies *progress* requirement
- When P_i exits CS, the selection of the next P_j to enter CS is arbitrary
 - No bounded waiting (it is a race!!!)

Operating Systems or Programming Language Support for Concurrency

- Solutions based on machine instructions such as *test and set* involve tricky coding
 - For example, the SetAndTest algorithm does not satisfy all the requirements to solve the critical-section problem
 - **Starvation** is possible
- We can build better solutions by providing synchronization mechanisms in the Operating System or Programming Language (This leaves the really tricky code to systems programmers)

Solution to bounded waiting

```
do{
    waiting[ i ] = TRUE;
    key = TRUE;
    while (waiting[ i ] && key)
        key = TestandSet(&lock);
    waiting[ i ] = FALSE;
    // Critical Section
    j=(i+1) % n ;
    while ((j != i) && !waiting[ j ])
        j=(j+1) % n ;
    if (j == i)
        lock = FALSE
    else
        waiting[ j ] = FALSE ;
    // remainder section
}while (TRUE);
```


Semaphores

- A Semaphore S is an integer variable that, apart from initialization, can only be accessed through 2 atomic and mutually exclusive operations:
 - wait(S)
 - sometimes called P()
 - Dutch proberen: “to test”
 - signal(S)
 - sometimes called V()
 - Dutch verhogen: “to increment”
- The classical definition of **wait** and **signal** is as shown in the following figures
- Useful when critical sections last for a short time, or we have lots of CPUs
- S initialized to positive value (to allow someone in at the beginning)

```
wait(S) {  
    while S<=0 do ;  
    S- -; }
```

```
signal(S) {  
    S++; }
```

Semaphores in Action


Initialize `mutex` to 1

Process P_i :

```
repeat
wait(mutex) ;
  CS
signal(mutex) ;
  RS
forever
```

Process P_j :

```
repeat
wait(mutex) ;
  CS
signal(mutex) ;
  RS
forever
```



```
wait(S) {
  while S <= 0 do ;
  S- -; }
```

```
signal(S) {
  S++; }
```

Synchronizing Processes using Semaphores

- Two processes:
 - P_1 and P_2
- Statement S_1 in P_1 needs to be performed **before** statement S_2 in P_2
- We want a way to make P_2 wait
 - Until P_1 tells it is OK to proceed

Define a semaphore “synch”
Initialize synch to 0

Put this in P_2 :
wait(synch);
 S_2 ;

And this in P_1 :
 S_1 ;
signal(synch);

Semaphores: the Problem of busy waiting

- Semaphore definitions (so far) all require busy waiting
- This type of semaphore is called **spinlock**
- This continual looping is a problem in a multiprogramming setting
- As a solution, modify the definition of the wait and signal semaphores

- Define a semaphore as a record

```
typedef struct {  
    int value;  
    struct process *L;  
} semaphore;
```

- Assume two simple operations:
 - **block** suspends the process that invokes it.
 - **wakeup(P)** resumes the execution of a blocked process **P**

Semaphore operations now defined as

```
wait(S):  
    S.value - -;  
    if (S.value < 0) {  
        add this process to S.L;  
        block; }
```

```
signal(S):  
    S.value++;  
    if (S.value <= 0) {  
        remove a process P from S.L;  
        wakeup(P); }
```

Deadlock and Starvation

- An implementation of a semaphore with a waiting queue may result in:
 - **Deadlock:** two or more processes are waiting indefinitely for an event that can be caused only by one of the waiting processes
 - Let S and Q be two semaphores initialized to 1

P0

```
wait(S);  
wait(Q);  
⋮  
signal(S);  
signal(Q)
```

P1

```
wait(Q);  
wait( S);  
⋮  
signal(Q);  
signal(S);
```

- **Starvation:** indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended
- If we add or remove processes from the list associated with a semaphore in LIFO manner

Applications of Semaphores

- Binary Semaphores
- Counting Semaphores
- Applications
 - Critical Section Problem
 - Deciding order of execution
 - For Managing Resources
 - e.g. 5 printers

Classical Problems of Synchronization

- Bounded-buffer problem
- Reader-Writer Problem
- Dining Philosophers Problem
- Monitors

```
wait(S) {  
    while S<=0 do ;  
    S- -; }
```

```
signal(S) {  
    S++; }
```

Classical Problems of Synchronization:

Bounded-Buffer Problem

- Shared data: **semaphore full**, **empty**, **mutex**;
- Initially: **full = 0**, **empty = n**, **mutex = 1**
- We have **n** buffers. Each buffer is capable of holding ONE item

Consumer

```
do {  
    wait(full)  
    wait(mutex);  
    ...  
    remove an item from  
    buffer to nextc  
    ...  
    signal(mutex);  
    signal(empty);  
    ...  
    consume the item in nextc  
    ...  
} while (1);
```

Producer

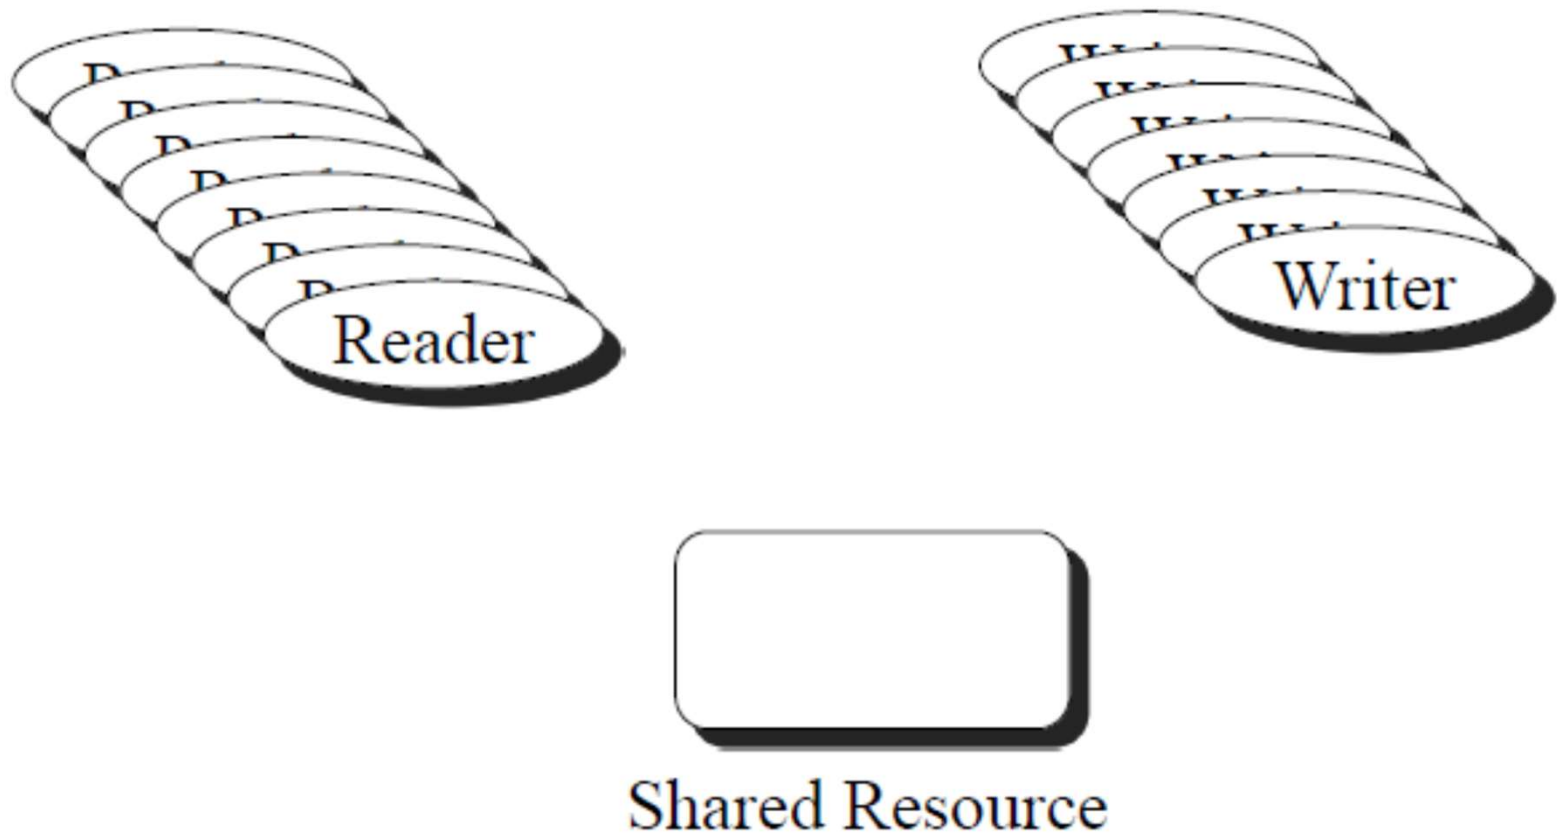
```
do {  
    ...  
    produce an item in  
    nextp  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    add nextp to buffer  
    ...  
    signal(mutex);  
    signal(full);  
} while (1);
```


Classical Problems of Synchronization:

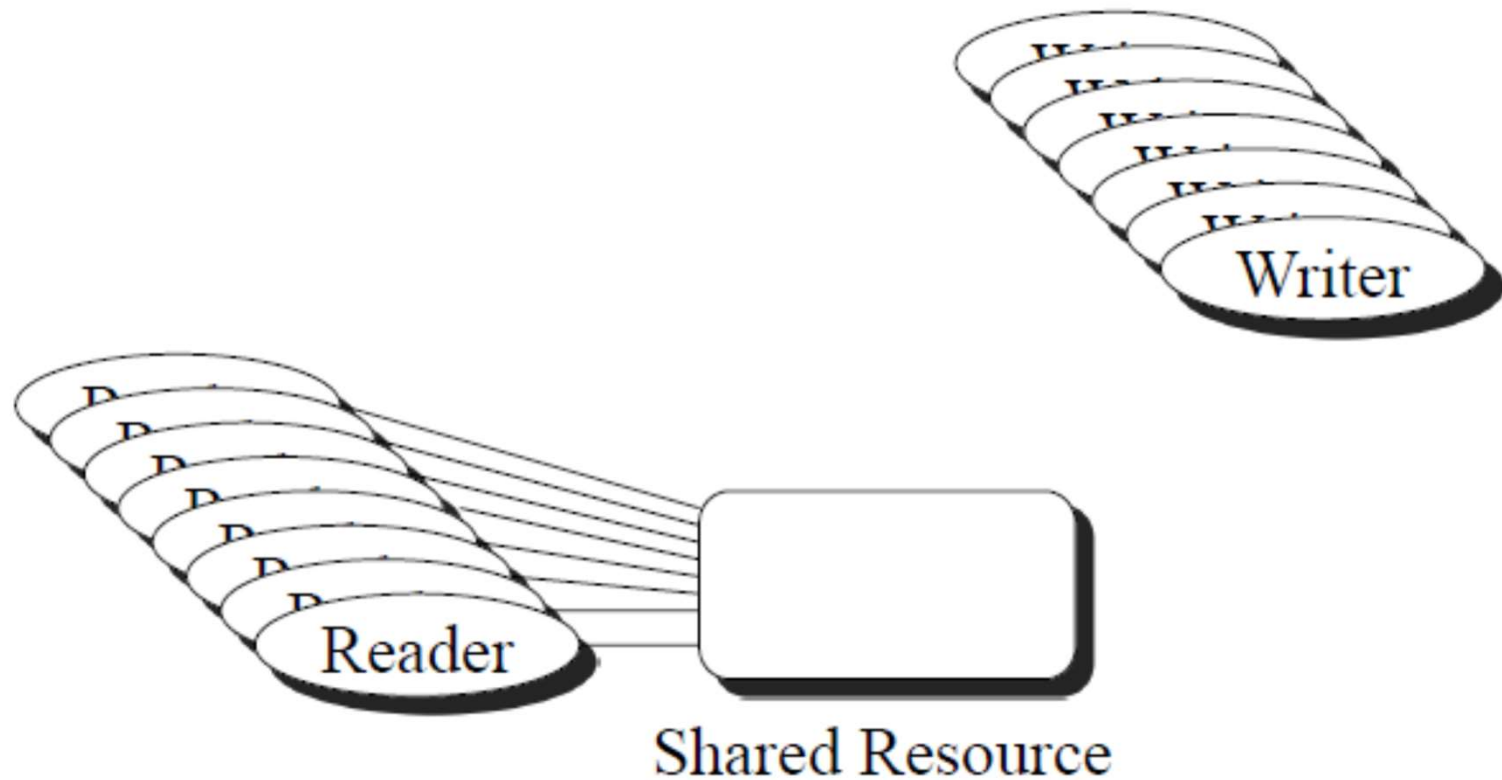
Readers-Writers Problem

- There is one writer and multiple readers
- The writer wants to write to the database
- The readers wants to read from the database
- We can not allow a writer and a reader writing and reading the database at the same time
- We can allow one or more readers reading from the database at the same time
- Two different versions:
 - **First reader-writers problem**
 - **Second readers-writers problem**

Reader-writers problem (Cont.)




Reader-writers problem (Cont.)






Concurrent readers

Reader-writers problem (Cont.)



First Solution: Reader's precedence

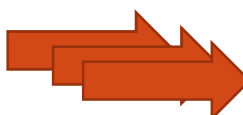
```
Reader() {  
    while(TRUE) {  
        wait(mutex);  
        readCount++;  
        if(readCount==1)  
             wait(wrt);  
        signal(mutex);  
  
        read(resource);  
  
        wait(mutex);  
        readCount--;  
        if(readCount == 0)  
             signal(wrt);  
        signal(mutex);  
    }  
}
```


```
Writer() {  
    while(TRUE) {  
         wait(wrt);  
        write(resource);  
        signal(wrt);  
    }  
}
```

```
resourceType *resource;  
int readCount = 0;  
semaphore mutex = 1;  
semaphore wrt = 1;
```

- First reader competes with writers
- Last reader signals writers

First Solution: reader's precedence

```
Reader() {  
    while(TRUE) {  
        wait(mutex);  
        readCount++;  
        if(readCount==1)  
            wait(wrt);  
        signal(mutex);  
        read(resource);  
        wait(mutex);  
        readCount--;  
        if(readCount == 0)  
            signal(wrt);  
        signal(mutex);  
    }  
}
```

```
Writer() {  
    while(TRUE) {  
        wait(wrt);  
        write(resource);  
        signal(wrt);  
    }  
}
```

- First reader competes with writers
- Last reader signals writers
- Any writer must wait for all readers
- Readers can starve writers
- “Updates” can be delayed forever

Second Solution: Writer's precedence

■ 2 ➔

```
Reader() {  
    while(TRUE) {  
        wait(rd);  
        wait(mutex1);  
        readCount++;  
        if(readCount == 1)  
            wait(wrt);  
        signal(mutex1);  
        signal(rd);  
        read(resource);  
        wait(mutex1);  
        readCount--;  
        if(readCount == 0)  
            signal(wrt);  
        signal(mutex1);  
    }  
}
```

```
writer() {  
    while(TRUE) {  
        wait(mutex2);  
        writeCount++;  
        if(writeCount == 1)  
            wait(rd);  
        signal(mutex2);  
        wait(wrt);  
        write(resource);  
        signal(wrt);  
        wait(mutex2);  
        writeCount--;  
        if(writeCount == 0)  
            signal(rd);  
        signal(mutex2);  
    }  
}
```

■ 1 ➔

■ 3 ➔

■ 4 ➔

int readCount = 0, writeCount = 0;
semaphore mutex1 = 1, mutex2 = 1;
semaphore rd = 1, wrt = 1;


The Dining Philosophers Problem

- A classical synchronization problem
- 5 philosophers who only eat and think
- Each need to use 2 forks for eating
- There are only 5 forks
- Illustrates the difficulty of allocating resources among process without deadlock and starvation



The Dining Philosopher Problem

- Each philosopher is a process
- One semaphore per fork:
 - Fork: array[0..4] of semaphores
 - Initialization:
fork[i].count:=1 for i:=0..4
- A first attempt:
 - Deadlock if each philosopher starts by picking his left fork!




```
Pi() {  
    while(TRUE){  
        think;  
        wait(fork[i]);  
        wait(fork[i+1 mod 5]);  
        eat;  
        signal(fork[i+1 mod 5]);  
        signal(fork[i]);  
    }  
}
```

The Dining Philosophers Problem

- Idea: admit only 4
- philosophers at a time who try to eat
- Then, one philosopher can always eat when the other 3 are holding one fork
- Solution: use another semaphore T to limit at 4 the number of philosophers “sitting at the table”
- Initialize: $T.count := 4$

```
Pi(){  
  while(TRUE){  
    think;  
    wait(T);  
    wait(fork[i]);  
    wait(fork[i+1 mod 5]);  
    eat;  
    signal(fork[i+1 mod 5]);  
    signal(fork[i]);  
    signal(T);  
  }  
}
```




Recall: Problems with Semaphores

- Semaphores are a powerful tool for enforcing mutual exclusion and coordinate processes
- **Problem:** wait(S) and signal(S) are scattered among several processes
 - It is difficult to understand their effects
 - Usage must be correct in all processes
 - One bad (or malicious) process can fail the entire collection of processes

Monitors

- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- Only one process may be active within the monitor at a time

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { .... }
    ...
    procedure Pn (...) { ..... }
    Initialization code ( .... ) { ... }
}
```



Monitors

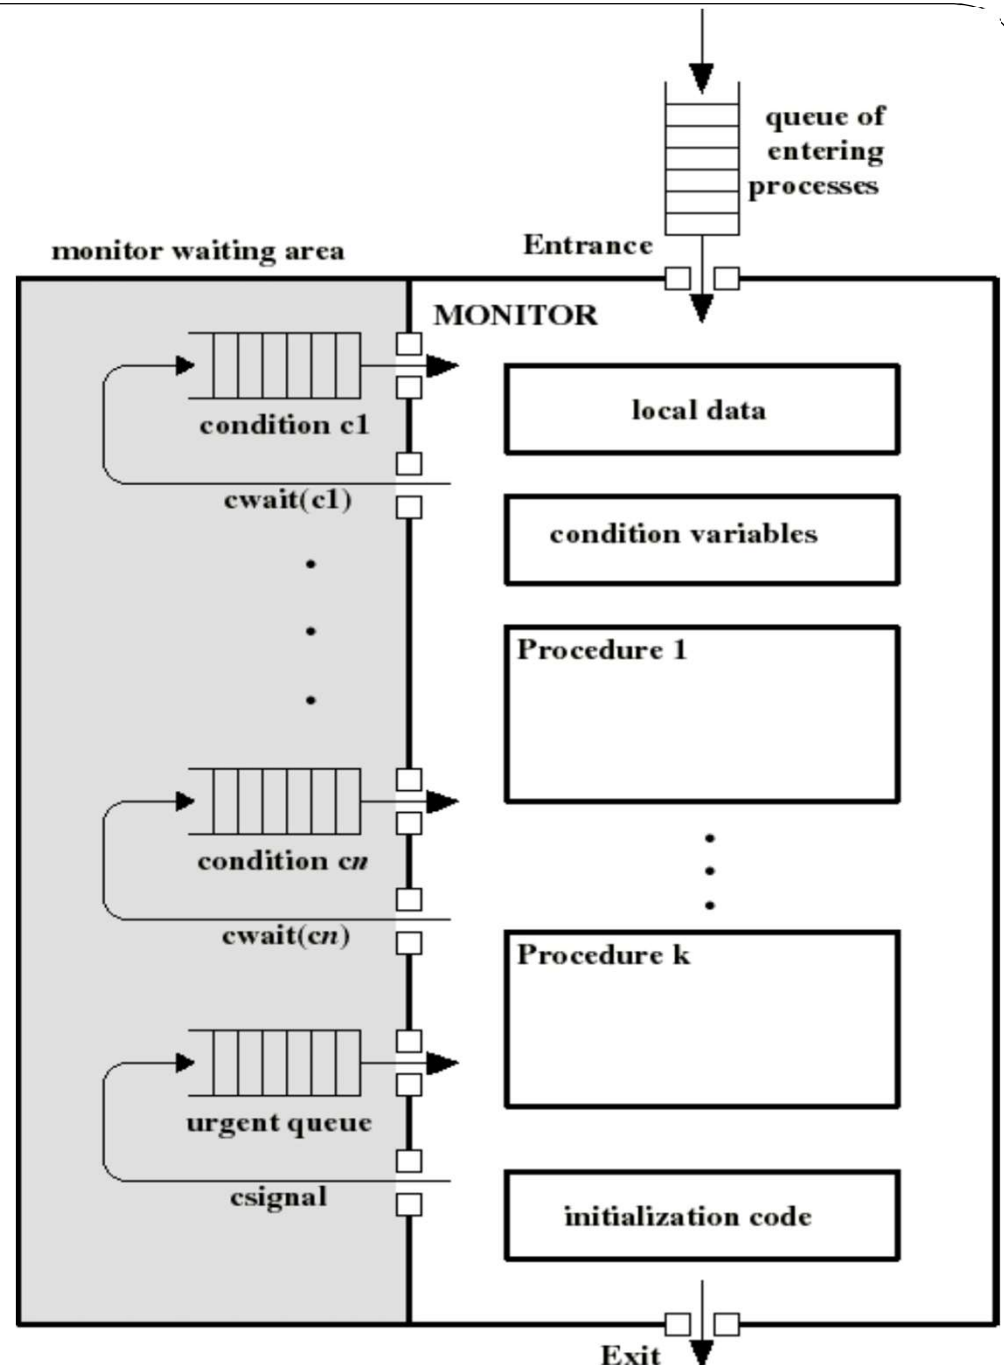
- Is a software module containing:
 - one or more procedures
 - an initialization sequence
 - local shared data variables
- Characteristics:
 - Local shared variables accessible only by monitor's procedures
 - a process enters the monitor by invoking one of its procedures
 - only one process can be in the monitor at any one time
- The monitor ensures mutual exclusion
 - no need to program this constraint explicitly
- Shared data are protected by placing them in the monitor
 - The monitor locks the shared data on process entry

Condition Variables

- Process synchronization is done using condition variables, which represent conditions a process may need to wait for before executing in the monitor
- `condition x, y;`
- Local to the monitor (accessible only within the monitor)
- Can be accessed and changed only by two functions:
 - `x.wait()`: blocks execution of the calling process on condition x
 - the process can resume execution only if another process executes `x.signal()`
 - `x.signal()`: resume execution of some process blocked on condition x.
 - If several such processes exists: choose any one
 - If no such process exists: do nothing

Monitors

- Awaiting processes are either in the entrance queue or in a condition queue
- A process puts itself into condition queue *cn* by issuing *cn.wait()*
- *cn.signal()* brings into the monitor one process in condition *cn* queue
- *signal-and-wait* and *signal-and-continue*



Producer Consumer using Monitor

- Two types of processes:
 - producers
 - consumers
- Synchronization is now confined within the monitor
- **append(.)** and **take(.)** are procedures within the monitor: are the only means by which P/C can access the buffer
- If these procedures are correct, synchronization will be correct for all participating processes

Producer:

```
while(TRUE){  
    produce item;  
    append(item);  
}
```

Consumer:

```
while(TRUE){  
    item=take();  
    consume item;  
}
```


Monitor for the Bounded P/C Problem

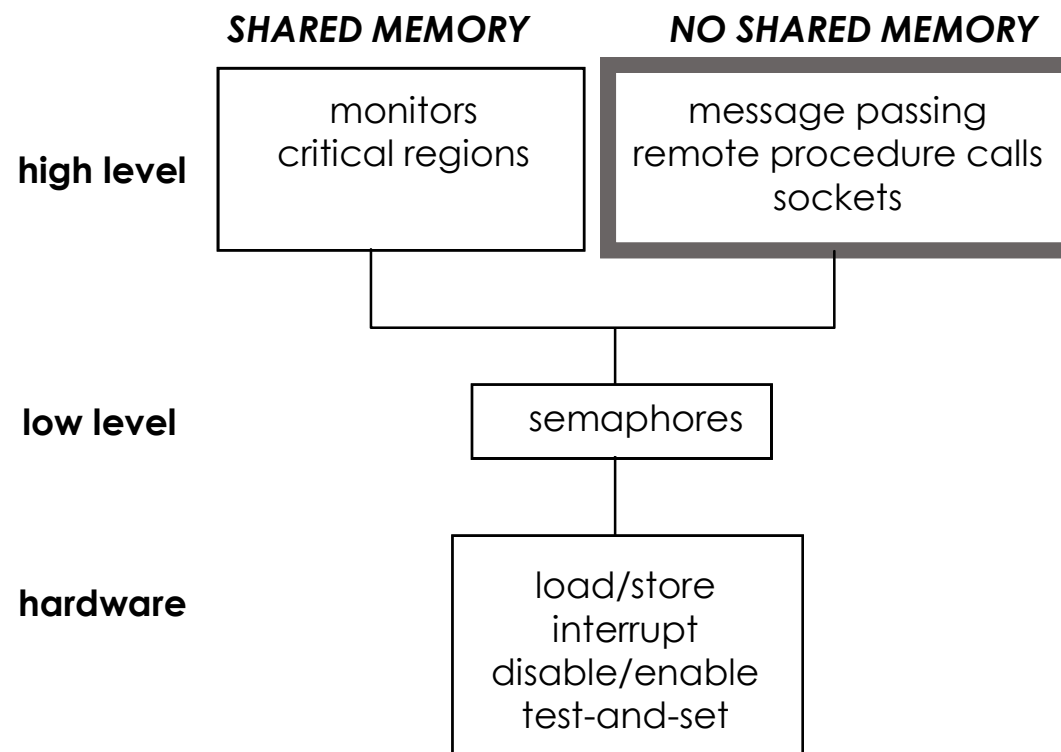
- Buffer:
 - *buffer*: array[0..k-1] of items;
- Buffer pointers and counts:
 - *nextin*: points to next item to be appended
 - *nextout*: points to next item to be taken
 - *count*: holds the number of items in the buffer
- Condition variables:
 - *notfull*: *notfull.signal()* indicates that the buffer is not full
 - *notempty*: *notempty.signal()* indicates that the buffer is not empty

Monitor for bounded buffer problem

```
Monitor boundedbuffer {  
    Item buffer[k];  
    integer nextin, nextout, count;  
    condition notfull, notempty;  
    Append(v){  
        if (count==k)  
            notfull.wait();  
        buffer[nextin] = v;  
        nextin = (nextin+1) mod k;  
        count++;  
        notempty.signal();  
    }  
    initialization_code(){  
        nextin=0; nextout=0; count=0;  
    }  
}
```

```
Item Take(){  
    if (count==0)  
        notempty.wait();  
    v = buffer[nextout];  
    nextout =  
        (nextout+1) mod k;  
    count--;  
    notfull.signal();  
    return v;  
}
```

Synchronization Primitives — Summary



References

- Operating System Concepts (Silberschatz, 8th edition)
Chapter 6